



DS 502 –Statistical Methods for Data Science

Assignment 2

Submitted By:

Siddhant Bhavsar and Jay Trivedi

Assignment 2

1. Using little bit of algebra, prove that (4,2) is equivalent to (4,3). In other words, the logistic function representation and logit representation for the logistic regression model are equivalent.

4.2 equation is $p(X) = (e^{\beta_0 + \beta_1 X}) / (1 + e^{\beta_0 + \beta_1 X})$

4.3 equation is $p(X) / (1 - p(X)) = e^{\beta_0 + \beta_1 X}$

$$p(X) = (e^{\beta_0 + \beta_1 X}) / (1 + e^{\beta_0 + \beta_1 X})$$

Subtract both side from 1,

$$1 - p(X) = 1 - (e^{\beta_0 + \beta_1 X}) / (1 + e^{\beta_0 + \beta_1 X})$$

$$1 - p(X) = 1 / (1 + e^{\beta_0 + \beta_1 X})$$

Now, taking reciprocal on both sides,

$$1 / (1 - p(X)) = 1 + e^{\beta_0 + \beta_1 X}$$

Now multiplying both sides by $p(X)$,

$$p(X) / (1 - p(X)) = ((e^{\beta_0 + \beta_1 X}) / (1 + e^{\beta_0 + \beta_1 X})) * (1 + e^{\beta_0 + \beta_1 X})$$

So,

$$p(X) / (1 - p(X)) = e^{\beta_0 + \beta_1 X}$$

Hence, Proved.

2. We now examine the differences between LDA and QDA.

a. If the Bayes decision boundary is linear, do we except LDA or QDA to perform better on the training set? On the test set?

QDA will perform better on training data as the Bayes decision boundary is linear and QDA is more flexible, whereas LDA will perform better than QDA on the test set. LDA by definition is

ought to have lower variance than QDA hence will perform better on test data as QDA might overfit based on the linear Bayes decision boundary.

- b. If the Bayes decision boundary is non-linear, do we expect LDA or QDA to perform better on the training set? On the test set?**

In this situation if the Bayes decision boundary is non-linear, then the QDA will perform better on both the set, i.e., training set as well as the test set. LDA usually perform bad on the non-linear Bayes decision boundary as it underfits the data and ends up having on more bias.

- c. In general, as the sample size n increases, do we expect the test prediction accuracy of QDA relative to LDA to improve, decline, or be unchanged? Why?**

We expect the test prediction accuracy of QDA relative to LDA to improve because QDA has high flexibility which tends to create a good fit. This is since adding more sample reduces variance which is good for QDA as problem often with QDA is variance and not bias thus QDA will end up giving better result.

- d. True or False: Even if the Bayes decision boundary for a given problem is linear, we will probably achieve a superior test error rate using QDA rather than LDA because QDA is flexible enough to model a linear decision boundary. Justify your answer.**

False. The variance by using a flexible method such as QDA leads to an inferior test error rate for fewer sample points. It might even cause overfit.

3. Suppose that we take a data set, divide it into equally-sized training and test sets, and then try out two different classification procedures. First, we use logistic regression and get an error rate of 20% on the training data and 30 % on the test data. Next, we use 1-nearest neighbors (i.e. $K = 1$) and get an average error rate (averaged over both test and training and test data sets) of 18%. Based on these results, which method should we prefer to use for classification of new observations? Why?

On basis of given data, we can infer that test error rate for $K=1$ is 36%. This is due to the fact that average error is 18% and for $K=1$ training error rate is 0% because for any training observation the class is the sample itself making test error rate 36%.

In case of logistic regression, we have lesser test error, thus we should prefer logistic regression in this case. Another reason for preferring logistic regression is for the fact that adding a new sample to the training data will change the whole decision boundary significantly, in other words the variance of this model is very high. Thus, it's better to use logistic regression model here.

4. This question should be answered using the weekly data set, which is part of the ISLR package. This data is similar in nature to the Smarket data from this chapter's lab, except that it contains 1,089 weekly returns for 21 years, from the beginning of 1990 to the end of 2010.

- a. **Produce some numerical and graphical summaries of the weekly data. Do there appear to be any patterns?**

Code:

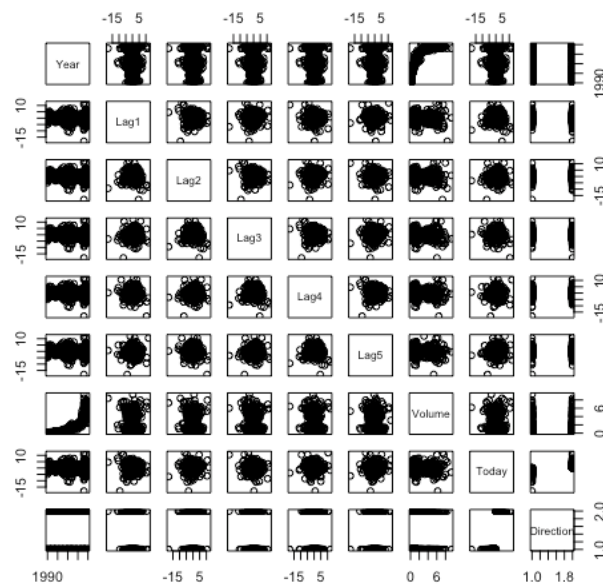
```
library(ISLR)
summary(Weekly)
head(Weekly)
pairs(Weekly)
plot(Weekly$Year,Weekly$Volume, xlab = 'Years', ylab = 'Volume', main = 'Plot of
Years vs Volume' )
```

Output:


```
> library(ISLR)
> summary(Weekly)
  Year      Lag1      Lag2      Lag3      Lag4      Lag5
Min. :1990 Min. :-18.1950 Min. :-18.1950 Min. :-18.1950 Min. :-18.1950 Min. :-18.1950
1st Qu.:1995 1st Qu.: -1.1540 1st Qu.: -1.1540 1st Qu.: -1.1580 1st Qu.: -1.1580 1st Qu.: -1.1660
Median :2000 Median : 0.2410 Median : 0.2410 Median : 0.2410 Median : 0.2380 Median : 0.2340
Mean :2000 Mean : 0.1506 Mean : 0.1511 Mean : 0.1472 Mean : 0.1458 Mean : 0.1399
3rd Qu.:2005 3rd Qu.: 1.4050 3rd Qu.: 1.4090 3rd Qu.: 1.4090 3rd Qu.: 1.4090 3rd Qu.: 1.4050
Max. :2010 Max. :12.0260 Max. :12.0260 Max. :12.0260 Max. :12.0260 Max. :12.0260
  Volume      Today      Direction
Min. :0.08747 Min. :-18.1950 Down:484
1st Qu.:0.33202 1st Qu.: -1.1540 Up :605
Median :1.00268 Median : 0.2410
Mean :1.57462 Mean : 0.1499
3rd Qu.:2.05373 3rd Qu.: 1.4050
Max. :9.32821 Max. :12.0260

> head(Weekly)
  Year Lag1 Lag2 Lag3 Lag4 Lag5 Volume. Today Direction
1 1990 0.816 1.572 -3.936 -0.229 -3.484 0.1549760 -0.270 Down
2 1990 -0.270 0.816 1.572 -3.936 -0.229 0.1485740 -2.576 Down
3 1990 -2.576 -0.270 0.816 1.572 -3.936 0.1598375 3.514 Up
4 1990 3.514 -2.576 -0.270 0.816 1.572 0.1616300 0.712 Up
5 1990 0.712 3.514 -2.576 -0.270 0.816 0.1537280 1.178 Up
6 1990 1.178 0.712 3.514 -2.576 -0.270 0.1544440 -1.372 Down
```

Plots:



Plot of Years vs Volume



From the plot we can conclude that the Year Vs Volume shows a logarithmic pattern.

- b. Use the full data set to perform a logistic regression with “Direction” as the response and the five lag variables plus “Volume” as predictors. Use the summary function to print the results. Do any of the predictors appear to be statistically significant? If so, which ones?

Code:

```
weeklyglm <- glm(Direction~Lag1+Lag2+Lag3+Lag4+Lag5+Volume,family = "binomial",
data=Weekly)
summary(weeklyglm)
```

Output:

```
> weeklyglm <- glm(Direction~Lag1+Lag2+Lag3+Lag4+Lag5+Volume,family =
"binomial", data=Weekly)
> summary(weeklyglm)
```

Call:

```
glm(formula = Direction ~ Lag1 + Lag2 + Lag3 + Lag4 + Lag5 +
Volume, family = "binomial", data = Weekly)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-1.6949	-1.2565	0.9913	1.0849	1.4579

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	0.26686	0.08593	3.106	0.0019 **
Lag1	-0.04127	0.02641	-1.563	0.1181
Lag2	0.05844	0.02686	2.175	0.0296 *
Lag3	-0.01606	0.02666	-0.602	0.5469
Lag4	-0.02779	0.02646	-1.050	0.2937
Lag5	-0.01447	0.02638	-0.549	0.5833
Volume	-0.02274	0.03690	-0.616	0.5377

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 1496.2 on 1088 degrees of freedom
Residual deviance: 1486.4 on 1082 degrees of freedom
AIC: 1500.4

Number of Fisher Scoring iterations: 4

Lag2 seems to be the only predictor which is statistically significant. Even its Estimate coefficient is 0.058 which indicates increase of lag2 value by 1 will affect the result by exponential term $e^{0.058}$.

- c. **Compute the confusion matrix and overall fraction of correct predictions. Explain what the confusion matrix is telling you about the types of mistakes made by logistic regression.**

Code:

```
wall <- predict(weeklyglm, type = 'response')
pred <- ifelse(wall > 0.5, "Up", "Down")
confusion_matrix <- table(pred, Weekly$Direction)
confusion_matrix
```

Output:

```
pred Down Up
Down  54  48
Up    430 557
```

From the Output, we can conclude that the percentage of true positive value is $(557/605) * 100 = 92.066\%$ which is quite high compared to the percentage of true negative values which is $(54/484) * 100 = 11.15\%$. That means, the changes of prediction of market going up is quite higher than the prediction of downfall in the market.

- d. **Now fit the logistic regression model using a training data period from 1990 to 2008, with Lag2 as the only predictor. Compute the confusion matrix and the overall fraction of correct predictions for the held-out data (that is, the data from 2009 and 2010).**

Code:

```
early <- Weekly[Weekly$Year<2009,]
later <- Weekly[Weekly$Year>2008,]
earlyglm <- glm(Direction~Lag2, family = "binomial", data= early)
summary(earlyglm)

laterpredict <- predict(earlyglm, newdata = later, type="response")
laterdirs <- Weekly$Direction[Weekly$Year>2008]
summary(laterpred)
laterpred <- rep("Down", 104)
laterpred[laterpredict>0.5] <- "Up"
table(laterpred, laterdirs)
```

Output:

```
> summary(earlyglm)

Call:
glm(formula = Direction ~ Lag2, family = "binomial", data = early)

Deviance Residuals:
```

```

Min    1Q    Median    3Q    Max
-1.536 -1.264  1.021   1.091  1.368

```

Coefficients:

```

              Estimate    Std. Error z value Pr(>|z|)
(Intercept)  0.20326     0.06428   3.162  0.00157 **
Lag2         0.05810     0.02870   2.024  0.04298 *
---

```

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

```

Null deviance: 1354.7 on 984 degrees of freedom
Residual deviance: 1350.5 on 983 degrees of freedom
AIC: 1354.5

```

Number of Fisher Scoring iterations: 4

Code:

```

laterpredict <- predict(earlyglm, newdata = later, type="response")
laterdirs <- Weekly$Direction[Weekly$Year>2008]
summary(laterpred)

```

Output:

```

      Length Class  Mode
      104   character character

```

The length of laterpred dataframe is 104, so we assign them as Down.

Code:

```

laterpred <- rep("Down", 104)
laterpred[laterpredict>0.5] <- "Up"
table(laterpred, laterdirs)

```

Output:

```

      Laterdirs
laterpred. Down Up
Down       9   5
Up        34  56

```

From this output, we can conclude that the percentage of true positive value is $(56/61) * 100 = 91.80\%$ which is quite high compared to the percentage of true negative values which is $(9/43) * 100 = 20.93\%$. That means, the changes of prediction of market going up is quite higher than the prediction of downfall in the market.

Total true correct predictions using logistic regression = $65/104 * 100 = 62.50\%$

e. Repeat d. using LDA

Code:

```
library(MASS)
earlylda <- lda(Direction~Lag2, family = "binomial", data= early)
laterpredictlda <- predict(earlylda, newdata = later, type="response")
laterdirslda <- Weekly$Direction[Weekly$Year>2008]
table(laterpredictlda$class, laterdirslda)
```

Output:

	laterdirslda	
	Down	Up
Down	9	5
Up	34	56

From this output, we can conclude that the percentage of true positive value is $(56/61) * 100 = 91.80\%$ which is quite high compared to the percentage of true negative values which is $(9/43) * 100 = 20.93\%$. That means, the changes of prediction of market going up is quite higher than the prediction of downfall in the market.

Total true correct predictions using LDA = $65/104 * 100 = 62.50\%$

f. Repeat d. using QDA

Code:

```
earlyqda <- qda(Direction~Lag2, family = "binomial", data= early)
laterpredictqda <- predict(earlyqda, newdata = later, type="response")
laterdirsqda <- Weekly$Direction[Weekly$Year>2008]
table(laterpredictqda$class, laterdirsqda)
```

Output:

	laterdirsqda	
	Down	Up
Down	0	0
Up	43	61

QDA implies all the data is going up. Thus we can say QDA has the worst error rate which is $43/104 * 100 = 41.35\%$

Total true correct predictions using QDA = $61/104 * 100 = 58.65\%$

g. Repeat d. using KNN and k = 1

Code:

```
library(class)
```

```

set.seed(1)
earlyset <- (Weekly$Year<2009)
laterset <- (Weekly$Year>2008)
earlylag2 <- Weekly[earlyset,"Lag2",drop=F]
laterlag2 <- Weekly[laterset,"Lag2",drop=F]
predknn <- knn(earlylag2, laterlag2, early$Direction, k = 1)
table(predknn, laterdirs)

```

Output:

	laterdirs	
predknn	Down	Up
Down	21	30
Up	22	31

We can state that the percentage of correct prediction in this case is $52/104 \times 100 = 50\%$. That means the error rate too is 50%. Through this model we can also say that the true prediction of market goes up is $31/61 \times 100 = 50.82\%$ right. Total true correct predictions using KNN = $52/104 \times 100 = 50\%$

h. **Which of these methods appear to provide best results on this data?**

LDA and Logistic regression gives us the best result rate which is 62.50%, followed by QDA which is 58.65% and at last KNN for $k=1$ which is 50%.

i. **Experiment with different combinations of predictors, including transformations and interactions, for each of the methods. Report the variables, method, and associated confusion matrix that provides the best results on the held-out data. Note that you should also experiment with values for K in the KNN classifier.**

Logistic Regression using Lag2+Lag3+lag4+Volume.

Code:

```

earlyglm2 <- glm(Direction~Lag2+Lag3+Lag4+Volume, family = "binomial", data=
early)
summary(earlyglm2)
laterpredict2 <- predict(earlyglm2, newdata = later, type="response")
laterdirs2 <- Weekly$Direction[Weekly$Year>2008]
laterpred2 <- rep("Down", 104)
laterpred2[laterpredict2>0.5] <- "Up"
table(laterpred2, laterdirs2)

```

Output:

	laterdirs2	
laterpred2	Down	Up
Down	26	34
Up	17	27

Here we can see that the true prediction percentage is $53/104 * 100 = 50.96\%$ for logistic regression using Lag2+Lag3+Lag4+Volume.

Performing LDA using Lag2+Lag3+Lag4+Volume

Code:

```
library(MASS)
earlylda2 <- lda(Direction~Lag2+Lag3+Lag4+Volume, family = "binomial", data=
early)
laterpredictlda2 <- predict(earlylda2, newdata = later, type="response")
laterdirsllda2 <- Weekly$Direction[Weekly$Year>2008]
table(laterpredictlda2$class, laterdirsllda2)
```

Output:

	laterdirsllda2	
	Down	Up
Down	26	35
Up	17	26

Here we can see that the true prediction percentage is $52/104 * 100 = 50\%$ for LDA using Lag2+Lag3+Lag4+Volume.

Performing QDA by interacting Lag3 by Lag4

Code:

```
earlyqda2 <- qda(Direction~Lag3:Lag4, family = "binomial", data= early)
laterpredictqda2 <- predict(earlyqda2, newdata = later, type="response")
laterdirsqda2 <- Weekly$Direction[Weekly$Year>2008]
table(laterpredictqda2$class, laterdirsqda2)
```

Output:

	laterdirsqda2	
	Down	Up
Down	0	0
Up	43	61

Here we can see that the true prediction percentage is $62/104 * 100 = 59.62\%$ for QDA using Lag3:Lag4.

Performing KNN for k=10

Code:

```
predknn2 <- knn(earlylag2, laterlag2, early$Direction, k = 10)
table(predknn2, laterdirs)
```

Output:

	laterdirs	
predknn2	Down	Up
Down	17	19
Up	26	42

Here we performed KNN with $k=10$ on Lag2, it comes out that the true prediction percentage in this case is $59/104 \times 100 = 56.73\%$.

Among, all the operations that we performed it seems that QDA using Lag3:Lag4 gave us best true prediction ratio when the market is going up which is 59.62% true prediction.

5. In this problem, you will develop a model to predict whether a given car gets high or low gas mileage based on the Auto data set.

- a. Create a binary variable, `mpg01`, that contains a 1 if mpg contains a value above its median, and a 0 if mpg contains a value below its median. You can compute the median using the `median()` function. Note you may find it helpful to use the `data.frame()` function to create a single data set containing both `mpg01` and the other Auto variables.

Code:

```
summary(Auto)
auto = Auto
auto$mpg01 <- ifelse(auto$mpg > median(auto$mpg),1,0)
auto$mpg01
head(auto)
auto
```

Output:

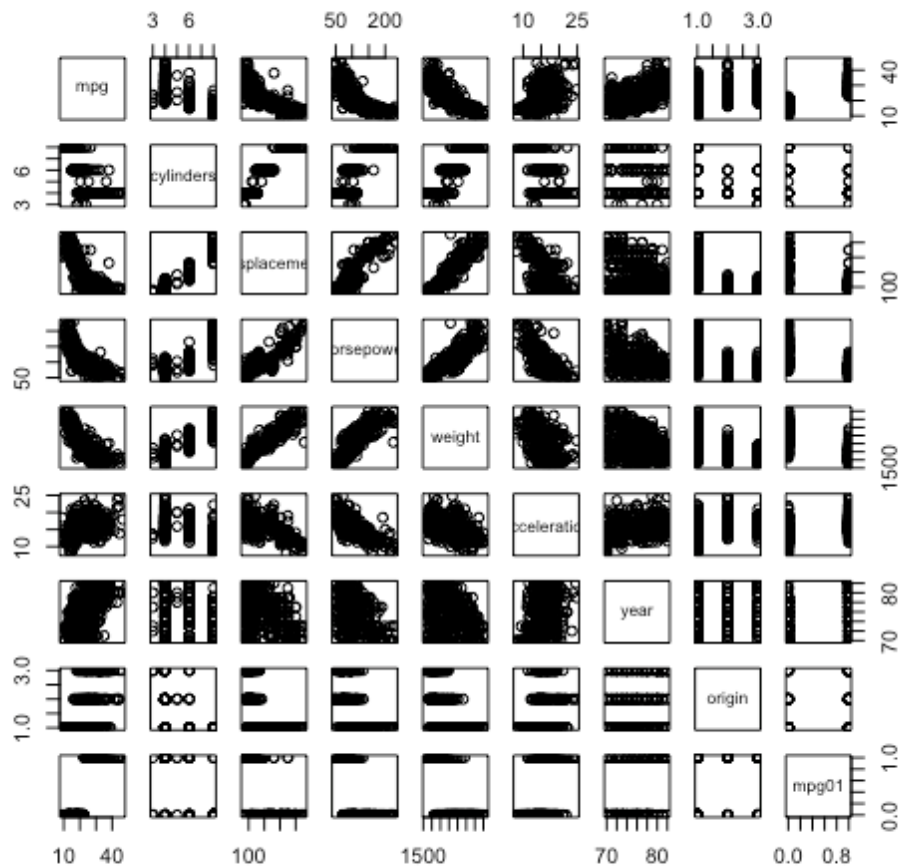
```
> head(auto)
  mpg cylinders displacement horsepower weight acceleration year origin      name      mpg01
1  18         8       307         130     3504         12.0     70    1 chevrolet chevelle malibu    0
2  15         8       350         165     3693         11.5     70    1   buick skylark 320      0
3  18         8       318         150     3436         11.0     70    1  plymouth satellite      0
4  16         8       304         150     3433         12.0     70    1    amc rebel sst      0
5  17         8       302         140     3449         10.5     70    1      ford torino      0
6  15         8       429         198     4341         10.0     70    1  ford galaxie 500      0
```

- b. Explore the data graphically in order to investigate the association between `mpg01` and the other features. Which of the other features seem most likely to be useful in predicting `mpg01`? Scatterplots and boxplots may be useful tools to answer this question. Describe your findings.

Code:

```
pairs(auto)
```

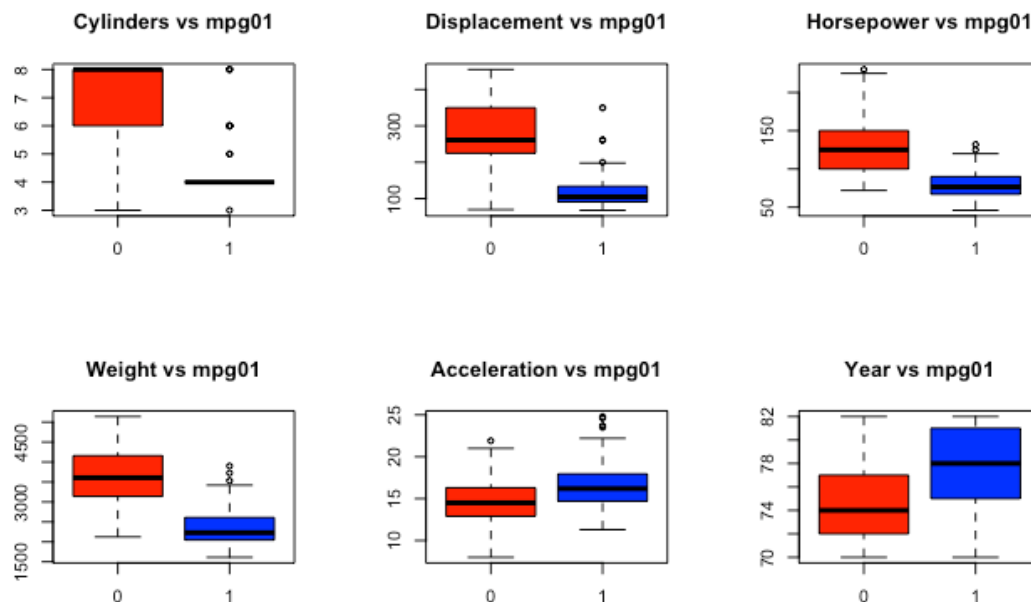
Plot:



Code:

```
par(mfrow=c(2,3))
boxplot(cylinders ~ mpg01, data = auto, main = "Cylinders vs mpg01", col=c("red",
"blue"))
boxplot(displacement ~ mpg01, data = auto, main = "Displacement vs mpg01",
col=c("red", "blue"))
boxplot(horsepower ~ mpg01, data = auto, main = "Horsepower vs mpg01",
col=c("red", "blue"))
boxplot(weight ~ mpg01, data = auto, main = "Weight vs mpg01", col=c("red",
"blue"))
boxplot(acceleration ~ mpg01, data = auto, main = "Acceleration vs mpg01",
col=c("red", "blue"))
boxplot(year ~ mpg01, data = auto, main = "Year vs mpg01", col=c("red", "blue"))
```

Plot:



Displacement, weight, cylinder, Horsepower have a strong correlation in predicting mpg01

c. Split the data into a training set and a test set.

Code:

```
set.seed(1)
a <- rnorm(nrow(auto))
summary(a)
test <- a > quantile(a,0.75)
train <- !test
trainset <- auto[train,]
testset <- auto[test,]
head(trainset)
head(testset)
```

Output:

```
> head(trainset)
  mpg cylinders displacement horsepower weight acceleration year origin      name mpg01
1  18         8       307      130 3504      12.0  70    1  chevrolet chevelle malibu  0
2  15         8       350      165 3693      11.5  70    1                buick skylark 320  0
3  18         8       318      150 3436      11.0  70    1      plymouth satellite  0
5  17         8       302      140 3449      10.5  70    1                ford torino  0
6  15         8       429      198 4341      10.0  70    1      ford galaxie 500  0
7  14         8       454      220 4354       9.0  70    1      chevrolet impala  0

> head(testset)
  mpg cylinders displacement horsepower weight acceleration year origin      name mpg01
4  16         8       304      150 3433      12.0  70    1                amc rebel sst  0
8  14         8       440      215 4312       8.5  70    1      plymouth fury iii  0
11 15         8       383      170 3563      10.0  70    1  dodge challenger se  0
```

15	24	4	113	95	2372	15.0	70	3	toyota corona mark ii	1
18	21	6	200	85	2587	16.0	70	1	ford maverick	0
19	27	4	97	88	2130	14.5	70	3	datson pl510	1

- d. **Perform LDA on the training data in order to predict mpg01 using the variables that seemed most associated with mpg01 in (b). What is the test error of the model obtained?**

Code:

```
lda1 <- lda(mpg01 ~ displacement+horsepower+weight+cylinders, data=trainset)
ldapred <- predict(lda1, testset)
table(testset$mpg01, ldapred$class)
mean(testset$mpg01 == ldapred$class)
```

Output:

```
> table(testset$mpg01, ldapred$class)

      0      1
0     36     8
1      3    51
> mean(testset$mpg01 == ldapred$class)
[1] 0.8877551
```

The test error of the model is $100 - 88.78 = 11.22\%$.

- e. **Perform QDA on the training data in order to predict mpg01 using the variables that seemed most associated with mpg01 in (b). What is the test error of the model obtained?**

Code:

```
qda1 <- qda(mpg01 ~ displacement+horsepower+weight+cylinders, data=trainset)
qdapred <- predict(qda1, testset)
table(testset$mpg01, qdapred$class)
mean(testset$mpg01 == qdapred$class)
```

Output:

```
> table(testset$mpg01, qdapred$class)

      0.      1
0     37     7
1      7    47
> mean(testset$mpg01 == qdapred$class)
[1] 0.8571429
```

The test error of this model is $100 - 85.71 = 14.29\%$

- f. **Perform logistic regression on the training data in order to predict mpg01 using the variables that seemed most associated with mpg01 in (b). What is the test error of the model obtained?**

Code:

```

earlyglm3 <- glm(mpg01~displacement+horsepower+weight+cylinders, family =
"binomial", data= trainset)
summary(earlyglm3)
laterpredict3 <- predict(earlyglm3, newdata = testset, type="response")
length(laterpredict3)
laterpred3 <- rep("0", 98)
laterpred3[laterpredict3>0.5] <- "1"
table(laterpred3, testset$mpg01)
mean(laterpred3 != testset$mpg01)

```

Output:

```

> table(laterpred3, testset$mpg01)

laterpred3 0 1
          0 36 6
          1  8 48
> mean(laterpred3 == testset$mpg01)
[1] 0.8571429
> mean(laterpred3 != testset$mpg01)
[1] 0.1428571

```

The test error rate of this model is 14.28%

- g. Perform KNN on the training data, with several values of K, in order to predict mpg01. Use only the variables that seemed most associated with mpg01 in (b). What test errors do you obtain? Which value of K performs the best on this data set?**

Code:

```

set.seed(1)
Autotrain = trainset[,c("displacement","horsepower","weight","acceleration")]
Autotest = testset[,c("displacement","horsepower","weight","acceleration")]
predknn=knn(Autotrain,Autotest,trainset$mpg01,k=1)
table(predknn,testset$mpg01)
mean(predknn!=testset$mpg01)

predknn=knn(Autotrain,Autotest,trainset$mpg01,k=2)
table(predknn,testset$mpg01)
mean(predknn!=testset$mpg01)

predknn=knn(Autotrain,Autotest,trainset$mpg01,k=3)
table(predknn,testset$mpg01)
mean(predknn!=testset$mpg01)

predknn=knn(Autotrain,Autotest,trainset$mpg01,k=4)

```



```

table(predknn,testset$mpg01)
mean(predknn!=testset$mpg01)

predknn=knn(Autotrain,Autotest,trainset$mpg01,k=5)
table(predknn,testset$mpg01)
mean(predknn!=testset$mpg01)

predknn=knn(Autotrain,Autotest,trainset$mpg01,k=6)
table(predknn,testset$mpg01)
mean(predknn!=testset$mpg01)

```

Output:

```

> set.seed(1)
> Autotrain = trainset[,c("displacement","horsepower","weight","acceleration")]
> Autotest = testset[,c("displacement","horsepower","weight","acceleration")]
> predknn=knn(Autotrain,Autotest,trainset$mpg01,k=1)
> table(predknn,testset$mpg01)

predknn 0 1
      0 35 5
      1 9 49
> mean(predknn!=testset$mpg01)
[1] 0.1428571
> predknn=knn(Autotrain,Autotest,trainset$mpg01,k=2)
> table(predknn,testset$mpg01)

predknn 0 1
      0 34 5
      1 10 49
> mean(predknn!=testset$mpg01)
[1] 0.1530612
> predknn=knn(Autotrain,Autotest,trainset$mpg01,k=3)
> table(predknn,testset$mpg01)

predknn 0 1
      0 36 4
      1 8 50
> mean(predknn!=testset$mpg01)
[1] 0.122449
> predknn=knn(Autotrain,Autotest,trainset$mpg01,k=4)
> table(predknn,testset$mpg01)

predknn 0 1

```

```
0 36 5
1 8 49
> mean(predknn!=testset$mpg01)
[1] 0.1326531
> predknn=knn(Autotrain,Autotest,trainset$mpg01,k=5)
> table(predknn,testset$mpg01)

predknn 0 1
0 36 7
1 8 47
> mean(predknn!=testset$mpg01)
[1] 0.1530612
> predknn=knn(Autotrain,Autotest,trainset$mpg01,k=6)
> table(predknn,testset$mpg01)

predknn 0 1
0 37 6
1 7 48
> mean(predknn!=testset$mpg01)
[1] 0.1326531
```

In this KNN model, test error rate is :

For k=1, Test Error Rate = 14.29%

For k=2, Test Error Rate = 15.30%

For k=3, Test Error Rate = 12.24%

For k=4, Test Error Rate = 13.26%

For k=5, Test Error Rate = 15.31%

For k=6, Test Error Rate = 13.27%

6. Using basic statistical properties of the variance, as well as single variable calculus, derive (5.6). In other words, prove that α given by (5.6) does indeed minimize $\text{Var}(\alpha X + (1 - \alpha)Y)$.

Hw 2 Question 6

To prove: $\alpha = \frac{\sigma_x^2 - \sigma_{xy}}{\sigma_x^2 + \sigma_y^2 - 2\sigma_{xy}}$

where $\sigma_x^2 = \text{Var}(X)$, $\sigma_y^2 = \text{Var}(Y)$ & $\sigma_{xy} = \text{Cov}(X, Y)$

Solution:

properties

$$1. \text{Var}(X+Y) = \text{Var}(X) + \text{Var}(Y) + 2\text{Cov}(X, Y)$$

$$2. \text{Var}(\alpha X) = \alpha^2 \text{Var}(X)$$

$$3. \text{Cov}(\alpha X, Y) = \alpha \text{Cov}(X, Y)$$

$$\text{We have } \text{Var}(\alpha X + (1-\alpha)Y) \\ = \text{Var}(\alpha X) + \text{Var}((1-\alpha)Y) + 2\text{Cov}(\alpha X, (1-\alpha)Y) \text{ by 1}$$

$$= \alpha^2 \text{Var}(X) + (1-\alpha)^2 \text{Var}(Y) + 2\alpha(1-\alpha) \text{Cov}(X, Y)$$

↳ by 2 & 3

$$= \alpha^2 \sigma_x^2 + (1-\alpha)^2 \sigma_y^2 + 2\alpha(1-\alpha) \sigma_{xy} \text{ by definition}$$

To minimize variance.

$$\frac{d}{d\alpha} \text{Var}(\alpha X + (1-\alpha)Y) = 0$$

$$\therefore 2\alpha \sigma_x^2 + 2\sigma_y^2(1-\alpha)(-1) + 2\sigma_{xy}(-2\alpha+1) = 0$$

$$2\alpha \sigma_x^2 + \sigma_y^2(\alpha-1) + \sigma_{xy}(-2\alpha+1) = 0$$

$$2(\sigma_x^2 + \sigma_y^2 - 2\sigma_{xy}) - \sigma_y^2 + \sigma_{xy} = 0$$

$$\alpha = \frac{\sigma_y^2 - \sigma_{xy}}{\sigma_x^2 + \sigma_y^2 - 2\sigma_{xy}}$$

Hence proved.

7. We will now derive the probability that a given observation is part of a bootstrap sample. Suppose that we obtain a bootstrap sample from a set of n observations.

(a) What is the probability that the first bootstrap observation is not the j th observation from the original sample? Justify your answer.

The probability that an observation is chosen is $1/n$ as every observation is equally like and the total probability is 1. Thus, the probability that a particular observation is not the j th observation is given by

$$1 - 1/n$$

(b) What is the probability that the second bootstrap observation is not the j th observation from the original sample?

The probability of second bootstrap observation is not the j th observation is still $(1-1/n)$ as bootstrap performs repeated sampling with replacement.

(c) Argue that the probability that the j th observation is not in the bootstrap sample is $(1 - 1/n)^n$.

As stated in questions above the probability of not choosing a bootstrap sample is $(1 - 1/n)$. Bootstrap by definition samples with replacement, hence new samples are independent of already chosen values.

So, when we are doing it n times the probability that j th observation is not in bootstrap sample simply gets multiplied. Hence the probability is given by:

$$(1-1/n) (1-1/n) (1-1/n) \dots n \text{ times} = (1 - 1/n)^n.$$

(d) When $n = 5$, what is the probability that the j th observation is in the bootstrap sample?

Probability that j th observation is in bootstrap sample is given by:

$$1 - (1 - 1/n)^n = 1 - (1 - 1/5)^5$$

$$= 0.672$$

(e) When $n = 100$, what is the probability that the j th observation is in the bootstrap sample?

Probability that j th observation is in bootstrap sample is given by:

$$1 - (1 - 1/n)^n = 1 - (1 - 1/100)^{100}$$

$$= 0.634$$

(f) When $n = 10,000$, what is the probability that the j th observation is in the bootstrap sample?

Probability that j th observation is in bootstrap sample is given by:

$$1 - (1 - 1/n)^n = 1 - (1 - 1/10000)^{10000}$$


$$= 0.632$$

(g) Create a plot that displays, for each integer value of n from 1 to 100,000, the probability that the j th observation is in the bootstrap sample. Comment on what you observe.

Code:

```
pr = function(n) return(1 - (1 - 1/n)^n)
x = 1:10000
plot(x, pr(x), col = 'blue', xlab = "Value of n", ylab = "Probability that jth is in bootstrap")
```

Output



The plot is exponentially decreasing and reaches value of 0.632 which perhaps acts as an asymptote.

(h) We will now investigate numerically the probability that a bootstrap sample of size $n = 100$ contains the j th observation. Here $j = 4$. We repeatedly create bootstrap samples, and each time we record whether or not the fourth observation is contained in the bootstrap sample.

```
> store=rep(NA, 10000) > for (i in 1:10000) {store[i]=sum (sample (1:100 , rep =TRUE)==4) >0}
```

```
} > mean(store)
```

Comment on the results obtained.

```
> set.seed(50)
> store=rep (NA, 10000)
> for (i in 1:10000)
+ store[i] <- sum (sample(1:100, rep=TRUE)==4) > 0
> mean(store)
[1] 0.6347
```

The code above samples 100 sample from 10000 and checks is 4 exists in the sample. As given by the plot in question before it should reach an asymptote of 0.632 as n tends to infinity limit of the function tends to 0.632.

Here as it still at 10000, thus limit has not been reached and the code gives value of 0.6347 for seed = 50.

Note seed has been set for reproducibility of the result

8. In Chapter 4, we used logistic regression to predict the probability of default using income and balance on the Default data set. We will now estimate the test error of this logistic regression model using the validation set approach. Do not forget to set a random seed before beginning your analysis.

(a) Fit a logistic regression model that uses income and balance to predict default.

Code

```
summary(Default)
attach(Default)
set.seed(50)
lr_m = glm(default~income+balance,family = binomial,data=Default)
summary(lr_m)
```

Output

```
Call:
glm(formula = default ~ income + balance, family = binomial,
    data = Default)

Deviance Residuals:
    Min       1Q   Median       3Q      Max
-2.4725 -0.1444 -0.0574 -0.0211  3.7245

Coefficients:
            Estimate Std. Error z value Pr(>|z|)
(Intercept) -1.154e+01  4.348e-01 -26.545 < 2e-16 ***
```

```
income 2.081e-05 4.985e-06 4.174 2.99e-05 ***
balance 5.647e-03 2.274e-04 24.836 < 2e-16 ***
---
Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 2920.6 on 9999 degrees of freedom
 Residual deviance: 1579.0 on 9997 degrees of freedom
 AIC: 1585

Number of Fisher Scoring iterations: 8

(b) Using the validation set approach, estimate the test error of this model. In order to do this, you must perform the following steps:

i. Split the sample set into a training set and a validation set.

Code

```
data <- sample(nrow(Default),nrow(Default)*0.8)
train <- Default[data,]
val <- Default[-data,]
summary(train)
summary(val)
```

Output

```
# train
default student balance income
No :7724 No :5635 Min. : 0.0 Min. : 772
Yes: 276 Yes:2365 1st Qu.: 477.9 1st Qu.:21270
      Median : 822.9 Median :34470
      Mean : 835.3 Mean :33478
      3rd Qu.:1167.3 3rd Qu.:43784
      Max. :2502.7 Max. :73554

# validation
default student balance income
No :1943 No :1421 Min. : 0.0 Min. : 2703
Yes: 57 Yes: 579 1st Qu.: 501.1 1st Qu.:21573
      Median : 825.1 Median :35054
      Mean : 835.5 Mean :33673
      3rd Qu.:1163.8 3rd Qu.:43891
      Max. :2654.3 Max. :70701
```

ii. Fit a multiple logistic regression model using only the training observations.

Code

```
lr_train <- glm(default~income+balance,family = binomial,data=train)
# summary of model
summary(lr_train)
```

Output

```
# summary of model
Call:
glm(formula = default ~ income + balance, family = binomial,
    data = train)

Deviance Residuals:
    Min       1Q   Median       3Q      Max
-2.5190 -0.1424 -0.0554 -0.0198  3.7343

Coefficients:
            Estimate Std. Error z value Pr(>|z|)
(Intercept) -1.167e+01  4.870e-01 -23.968 < 2e-16 ***
income       2.045e-05  5.470e-06   3.738 0.000185 ***
balance      5.761e-03  2.571e-04  22.407 < 2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

    Null deviance: 2400.8  on 7999  degrees of freedom
Residual deviance: 1279.1  on 7997  degrees of freedom
AIC: 1285.1
Number of Fisher Scoring iterations: 8
```

iii. Obtain a prediction of default status for each individual in the validation set by computing the posterior probability of default for that individual and classifying the individual to the default category if the posterior probability is greater than 0.5.

Code

```
#posterior probability
prob <- predict(lr_train,val,type = "response")
result <- ifelse(prob > 0.5,"Yes","No")
# confusion matrix
table(val$default,result,dnn=c("Actual","Predicted"))
```


Output

	Predicted	
Actual	No	Yes
No	1935	8
Yes	36	21

iv. Compute the validation set error, which is the fraction of the observations in the validation set that are misclassified.

Code

```
# error
round(mean(val$default!=result),digits = 3)
```

Output

```
[1] 0.022
```

(c) Repeat the process in (b) three times, using three different splits of the observations into a training set and a validation set. Comment on the results obtained.

Code

```
set.seed(50)

##### First iteration 20-80
data <- sample(nrow(Default),nrow(Default)*0.2)
train <- Default[data,]
val <- Default[-data,]
lr_train <- glm(default~income+balance,family = binomial,data=train)
prob <- predict(lr_train,val,type = "response")
result <- ifelse(prob > 0.5,"Yes","No")
round(mean(val$default!=result),digits = 3)

##### Second Iteration 50-50
data <- sample(nrow(Default),nrow(Default)*0.5)
train <- Default[data,]
val <- Default[-data,]
lr_train <- glm(default~income+balance,family = binomial,data=train)
prob <- predict(lr_train,val,type = "response")
result <- ifelse(prob > 0.5,"Yes","No")
round(mean(val$default!=result),digits = 3)
```

```
##### Third Iteration 80-20
data <- sample(nrow(Default),nrow(Default)*0.8)
train <- Default[data,]
val <- Default[-data,]
lr_train <- glm(default~income+balance,family = binomial,data=train)
prob <- predict(lr_train,val,type = "response")
result <- ifelse(prob > 0.5,"Yes","No")
round(mean(val$default!=result),digits = 3)
```

Output

```
#### First Iteration error
[1] 0.026

#### Second Iteration error
[1] 0.024

#### Third Iteration error
[1] 0.022
```

As we observe from the iterations the more data is given to train the model the error rate decreases, this is not always true as there is always a chance that a random validation data will not show this characteristic. It so turns out in our iterations that difference between error is constant, but it is not the case always. This is probably due to random seed that has been set.

(d) Now consider a logistic regression model that predicts the probability of default using income, balance, and a dummy variable for student. Estimate the test error for this model using the validation set approach. Comment on whether or not including a dummy variable for student leads to a reduction in the test error rate.

Code

```
set.seed(50)
data <- sample(nrow(Default),nrow(Default)*0.8)
train <- Default[data,]
val <- Default[-data,]
lr_train <- glm(default~income+balance+student,family = binomial,data=train)
prob <- predict(lr_train,val,type = "response")
result <- ifelse(prob > 0.5,"Yes","No")
round(mean(val$default!=result),digits = 3)

table(val$default,result,dnn=c("Actual","Predicted"))
```

Output

```
### error
[1] 0.021

#### Confusion matrix

      Predicted
Actual  No  Yes
No  1936   7
Yes   35  22
```

As we see from the result the error rate slightly decreases with the addition of student variable. As we can see in the confusion matrix this is due to increase of 1 in both true positive and true negative. The error decrease is very small so perhaps adding the student variable doesn't give significantly more information to the model.

9. We continue to consider the use of a logistic regression model to predict the probability of default using income and balance on the Default data set. In particular, we will now compute estimates for the standard errors of the income and balance logistic regression coefficients in two different ways: (1) using the bootstrap, and (2) using the standard formula for computing the standard errors in the glm() function. Do not forget to set a random seed before beginning your analysis.

(a) Using the summary() and glm() functions, determine the estimated standard errors for the coefficients associated with income and balance in a multiple logistic regression model that uses both predictors.

Code

```
set.seed(50)
####
lr_train <- glm(default~income+balance,family = binomial, data=Default)
# summary of model and standard error
summary(lr_train)$coef[,2]
```

Output

```
#### error
(Intercept)  income  balance
4.347564e-01 4.985167e-06 2.273731e-04
```

Thus, standard error coefficients associated with income and balance are 4.985167e-06 and 2.273731e-04 respectively.

(b) Write a function, `boot.fn()`, that takes as input the Default data set as well as an index of the observations, and that outputs the coefficient estimates for income and balance in the multiple logistic regression model.

Code

```
boot.fn = function(data, index) return(coef(glm(default~income+balance,family = binomial,data=
data, subset = index)))
```

(c) Use the `boot()` function together with your `boot.fn()` function to estimate the standard errors of the logistic regression coefficients for income and balance.

Code

```
library(boot)
boot(Default, boot.fn, 100)
```

Output

```
ORDINARY NONPARAMETRIC BOOTSTRAP

Call:
boot(data = Default, statistic = boot.fn, R = 100)

Bootstrap Statistics :
      original    bias  std. error
t1* -1.154047e+01 -1.431166e-02 4.963704e-01
t2*  2.080898e-05 -3.856758e-07 5.193719e-06
t3*  5.647103e-03  1.287655e-05 2.698000e-04
```

Thus, standard errors of the logistic regression coefficients for income and balance are 5.193719e-06 and 2.698000e-04 respectively.

(d) Comment on the estimated standard errors obtained using the `glm()` function and using your bootstrap function.

The answer from using bootstrap and glm function are similar but off by a bit around 2e-07 for income and 5e-05 for balance, but this will not be a problem as the variation is very small.